

Data Mining and Machine Learning

Clustering I

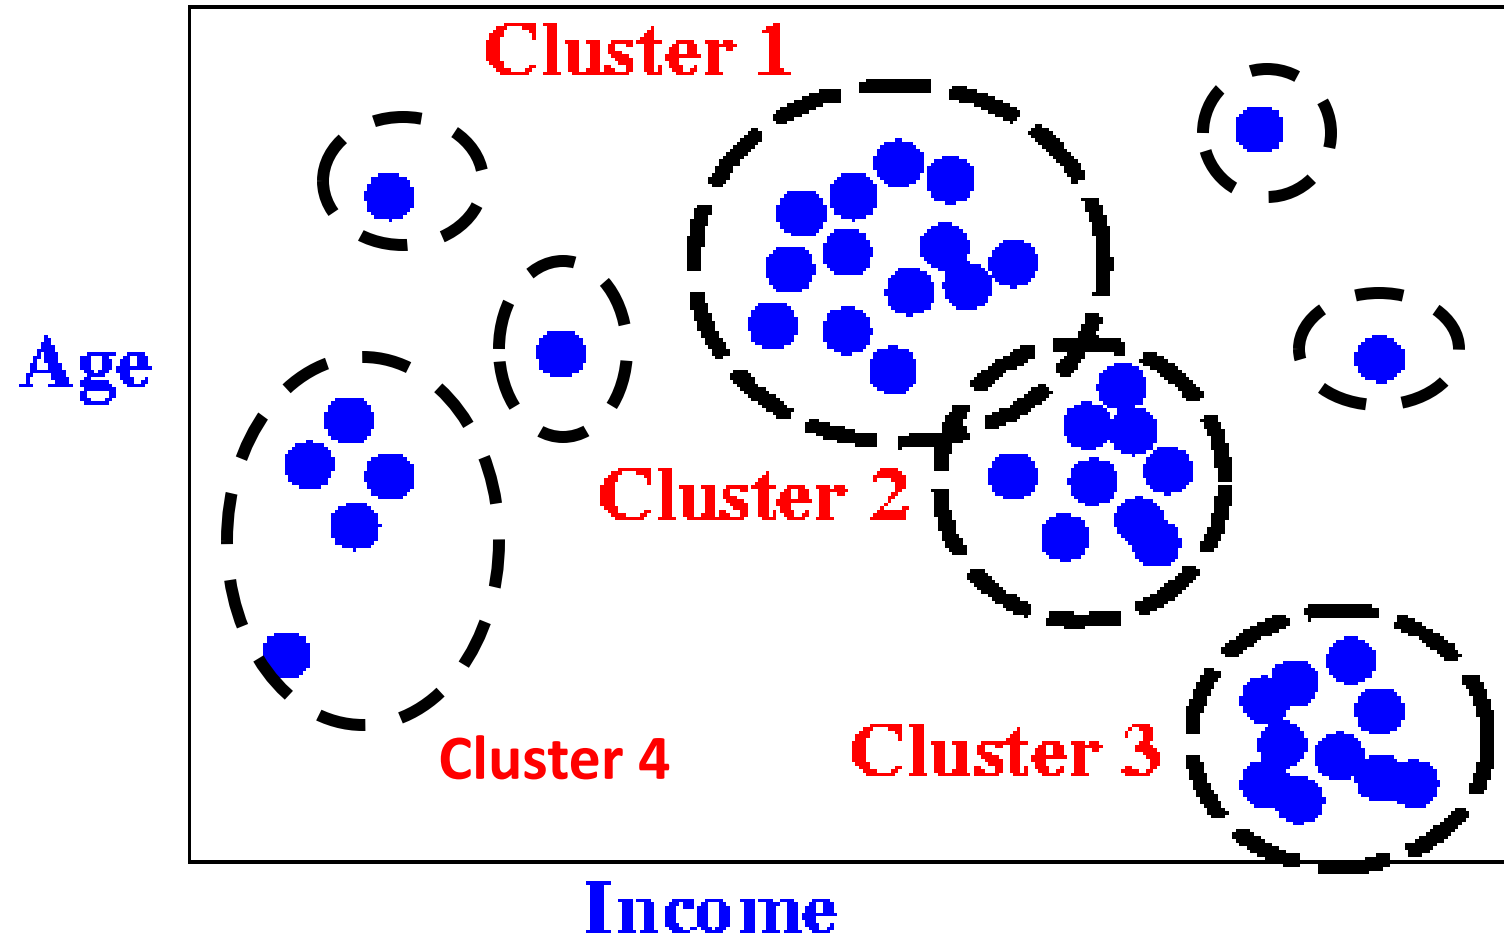
Dr. Edgar Acuna
Departamento de Matematicas

Universidad de Puerto Rico- Mayaguez

academic.uprm.edu/eacuna

Introduction

- The main idea of clustering is to group samples (rows) or features (columns) or both at the same time, according to the separation between them. The separation is determined by a distance measure called **dissimilarity measure**. It is assumed that the classes where the instances belong are unknown.
- For each sample, there exists a measurements vector $\mathbf{X}=(X_1,\dots X_G)$, here G is the number of features.
- The goal is to identify groups of similar samples based on n observed measurements $\mathbf{X}_1=\mathbf{x}_1,\dots,\mathbf{X}_n=\mathbf{x}_n$. For instance, if the \mathbf{x} 's represent levels of genetic expression obtained in microarrays of cancerous tumors, one could be able to identify the genetic characteristics of persons having distinct type of cancerous tumors.
- Clustering is also known as unsupervised classification or segmentation.



Introduction

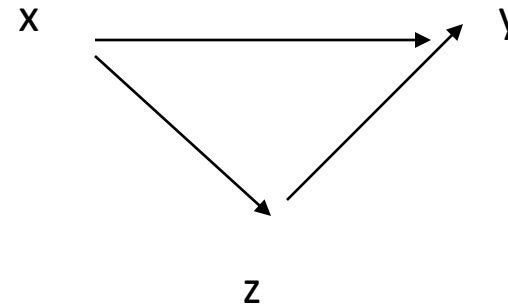
- When the number of columns is very large, these can be grouped according to their similarity and in this way the data dimensionality can be reduced. After that it is much easier to build a prediction model, since it is more convenient work with a small number of predictors.
- Also, clustering can be applied simultaneously to rows and columns of a dataset (Bi-clustering) (see Alon, et al, 1999, Getz et al , 2000 and, Lazaeronni and Owen, 2000)

Issues to take in account in clustering

- i) Which features can be used to form the groups? ,
- ii) Which dissimilarity measure should be used?.
- iii) Which clustering algorithm must be used?.
- iv) How many clusters should be formed?
- v) How to assign row (or columns) into clusters?
- vi) How to validate the formed clusters?

Properties of dissimilarity measures

- Non-negativity: $d(x,y) \geq 0$
- The distance of instance to itself is 0, $d(x,x) = 0$
- Symmetry: $d(x,y) = d(y,x)$
- Triangular Inequality:
 $d(x,y) \leq d(x,z) + d(z,y)$



Dissimilarity measures (continuous variables)

a) Minkowski Distance or L_p norm.

Particular cases:

$$D_p(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^M (x_i - y_i)^p \right)^{1/p}$$

- Euclidean Distance: $p=2$,

$$D_2 = \sqrt{\sum_{i=1}^M (x_i - y_i)^2}$$

- Manhattan or City-Block Distance:

$$D_1 = \sum_{i=1}^M |x_i - y_i|$$

- Chebychev Distance, $p=\infty$,

$$D_\infty = \max_{1 \leq i \leq M} |x_i - y_i|$$

Weighted Minkowski Distance:

$$D_p(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^M w_i (x_i - y_i)^p \right)^{1/p}$$

Dissimilarity measures (Cont)

b) Distances based on Quadratic forms.

$Q=(Q_{ij})$ is a definite positive square matrix of order $M \times M$ containing weights, then the quadratic distance between x and y is given by:

$$DQ(x, y) = [(x - y)' Q (x - y)]^{1/2} = \sqrt{\sum_{i=1}^M \sum_{j=1}^M (x_i - y_i) Q_{ij} (x_j - y_j)}$$

When $Q=V^{-1}$, where V is the covariance matrix between x and y , then DQ is called the Mahalanobis distance.

c) Canberra Distance.

$$D_{Can}(x, y) = \sum_{i=1}^M \frac{|x_i - y_i|}{|x_i + y_i|}$$

when x_i and y_i are both zeros then the i -th term of the sum is considered as zero.

Dissimilarity measures(nominal variables)

Hamming Distance. Let \mathbf{x} and \mathbf{y} be two vectors with the same dimension and taking values in $\Omega=\{0,1,2,\dots,k-1\}$, then the Hamming Distance (D_H) between them is defined as the number of different entries at the same position. Thus, if $\mathbf{x}=(1,1,0,0)$ and $\mathbf{y}=(0,1,0,1)$ then $D_H=2$.

- Also, it can be used for non-binary vectors. For example, if $\mathbf{x}=(0,1,3,2,1,0,1)$ and $\mathbf{y}=(1,1,2,2,3,0,2)$ then $DH(\mathbf{x},\mathbf{y})=4$.
- In case of binary vectors, the Hamming distance, the L2 distance and the L1 distance are the same.

Similarity measures (continuous variables))

Similarity=1-dissimilarity

a) The correlación measure:

$$r(x, y) = \frac{\sum_{i=1}^M (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^M (x_i - \bar{x})^2 \sum_{i=1}^M (y_i - \bar{y})^2}} = \frac{(x - \bar{x})'(y - \bar{y})}{\|x - \bar{x}\| \|y - \bar{y}\|}$$

Note: $1-s(x,y)$ can be considered as a dissimilarity measure. The cosine distance is similar to the correlation distance but in this case the sample means for x and y are taken as zeros.

b) The Tanimoto measure. It is defined by

$$S_T(x, y) = \frac{x' y}{\|x\|^2 + \|y\|^2 - x' y}$$

Also, it can be used with nominal variables.

Similarity measures (nominal variables)

i) The Tanimoto measure. Let X and Y be two vectors of length n_X and n_Y respectively. Let $n_{X \cap Y}$ be the cardinality of the intersection of X and Y , then the Tanimoto measure is defined by:

$$\frac{n_{X \cap Y}}{n_X + n_Y - n_{X \cap Y}}$$

This, the Tanimoto measure is the ratio between the number of elements in common to the two vector, divided by the number of elements of the union. In the case of binary variables, the Tanimoto becomes

$$\frac{a + d}{a + 2(c + b) + d}$$

where a represents the number of positions where both X and Y have the same value 0, d represents the number of coincidences where both X and Y take the value 1. While c and b represent the number of non-coincidences.

Similarity measures (nominal variables)

ii) **The coefficient of simple coincidences.** Defined by

$$\frac{a+d}{a+b+c+d}$$

iii) **The Jaccard-Tanimoto measure.** Defined by

$$\frac{d}{b+c+d}$$

iv) **The Russel –Rao measure.**

$$\frac{d}{a+b+c+d}$$

v) **The Dice-Czekanowski measure.** Similar to Jaccard-Tanimoto measure but assign double weight to the coincidences. That is,

$$\frac{2d}{b+c+2d}$$

Libraries and functions to compute distances in Python

The function `pairwise_distances` of `sklearn.metrics.pairwise` computes several distances including: `['cityblock', 'cosine', 'euclidean', 'l1', 'l2', 'manhattan']`, and from `scipy.spatial.distance`: `['braycurtis', 'canberra', 'chebyshev', 'correlation', 'dice', 'hamming', 'jaccard', 'kulsinski', 'mahalanobis', 'minkowski', 'rogerstanimoto', 'russellrao', 'seuclidean', 'sokalmichener', 'sokalsneath', 'sqeuclidean', 'yule']`

Example: Distances on Diabetes

```
url= "http://academic.uprm.edu/eacuna/diabetes.dat"
names = ['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class']
data = pd.read_table(url, names=names)
X=data.iloc[:,0:8]
#Normalizing the predictors
scaler = StandardScaler()
scaler.fit(X)
X= scaler.transform(X)
#Computing the Euclidean Distance
dist2=pairwise_distances(X)
```

Libraries and functions to compute distances in Python

Manhattan distance, also called CityBlock

```
dist_maha=pairwise_distances(X,metric="cityblock")
```

```
dist_maha
```

#Mahalanobis Distance

```
pairwise_distances(X,metric='mahalanobis',V=np.cov(X))
```

#Correlation Distance but between rows of Diabetes

```
distcor1=pairwise_distances(X,metric="correlation")
```

```
Distcor1
```

The class `sklearn.neighbors` has a function `DistanceMetric` for computing several distances

#Euclidean distance as a particular case of the Minkowski distance

```
dist_e=DistanceMetric.get_metric("minkowski",p=2)
```

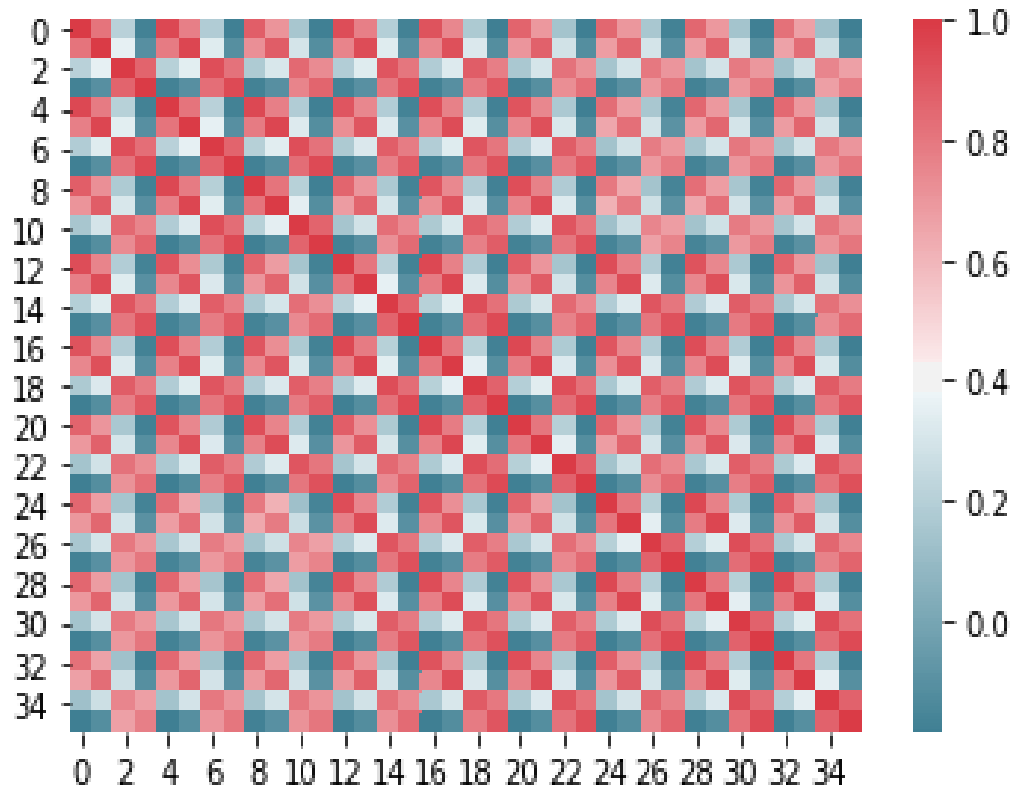
```
dist_e.pairwise(X)
```

Heatmaps

Heatmaps are useful to visualize the values of a distance metric in a dataset, since it is not easier to appreciate them in the numerical matrix.

Heatmap can be obtained using either matplotlib or seaborn

Here, we show the heatmap of the correlations among the features of landsat



Types of Clustering Algorithms

- I. Methods based on Partitioning (Kmeans, PAM,SOM)
- II. Hierarchical Methods.
- III. Methods based on density (DBSCAN): It groups in the same cluster the instances that have a lot of nearest neighbors. Outliers are identified.
- IV. Methods based on Models (Gaussian Mixtures)
- V. Methods based on Grids (STING, CLIQUE)

I. Partitioning Methods

The dataset is partitioned in a pre-specified number K of clusters, and then iteratively the data points are re-assigned to a cluster until a stopping criterion is satisfied. Usually until a function such as the sum of squares within each clusters is minimized.

Examples: K-means, PAM, CLARA, SOM, Clustering based on Gaussian Mixtures, Fuzzy K-means etc.

1-The k-means Algorithm (MacQueen, 1967).

The goal is to minimize the dissimilarity of the data points within each cluster and maximize the dissimilarity of the data points belonging to distinct clusters.

INPUT: A dataset S and the number of clusters k to be formed;

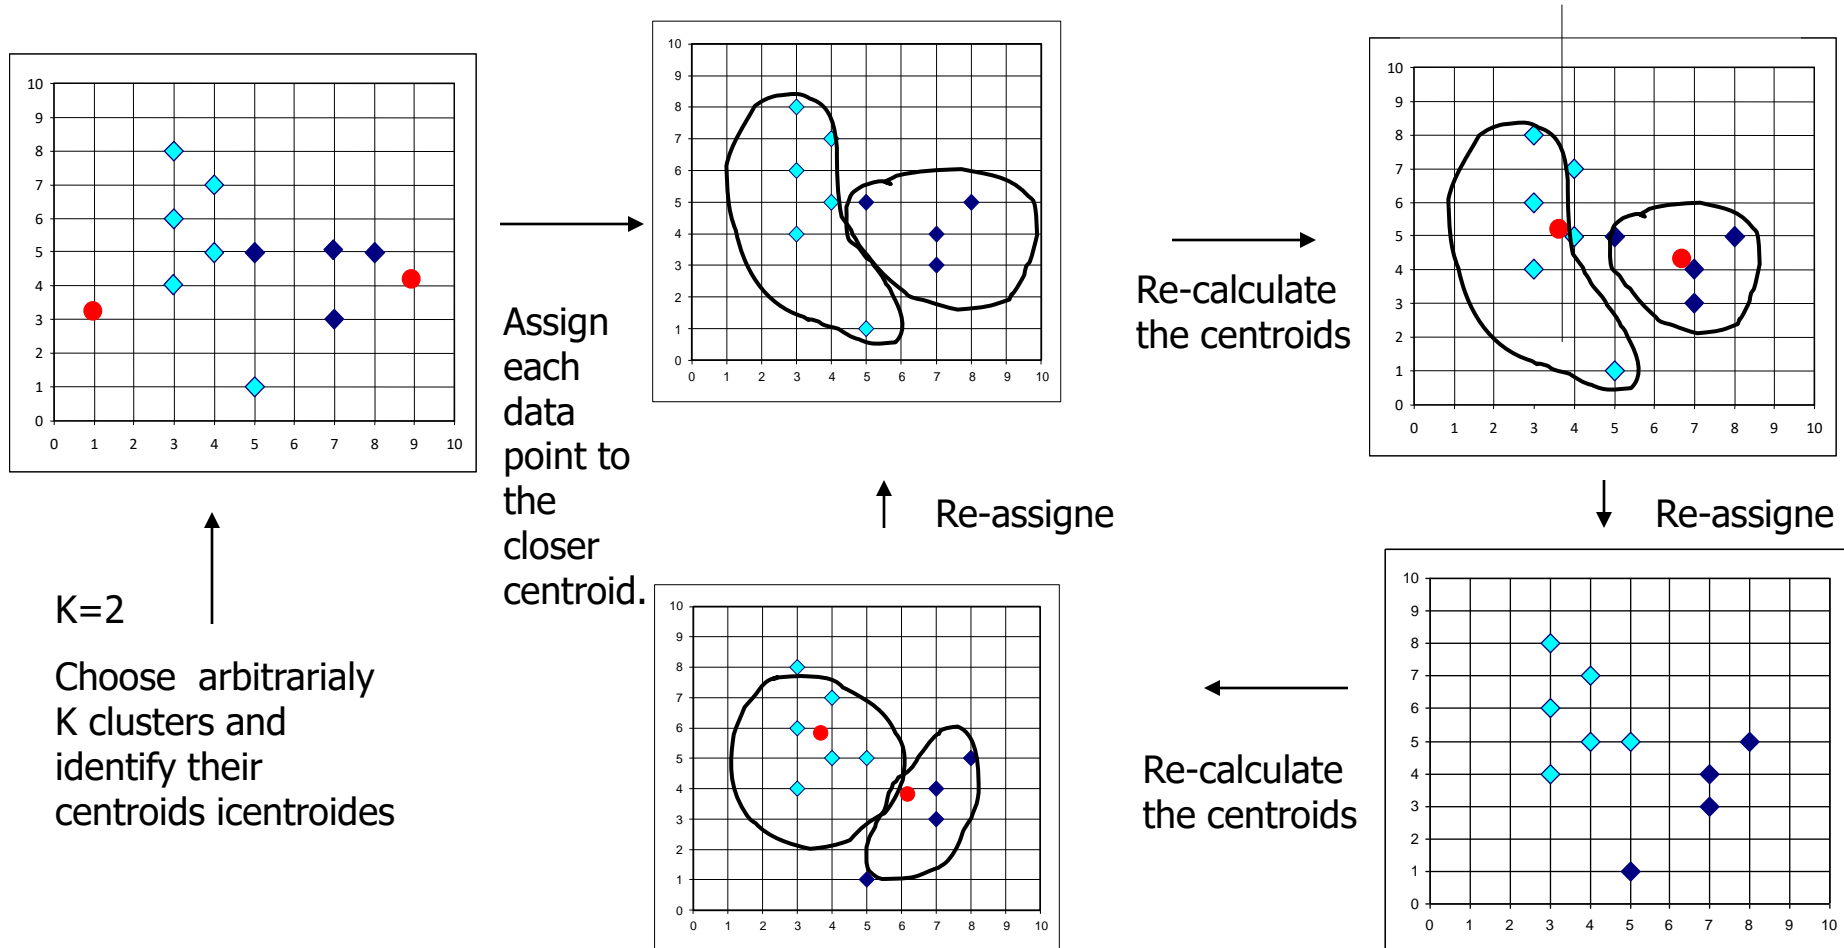
OUTPUT: A list L of the clusters where each data point of S is assigned.

1. Select the initial centroids of the K clusters: c_1, c_2, \dots, c_K .
2. Assign each data point x_i of S to the cluster C_i , which centroid c_i is closer to x_i . That is, $C_i = \operatorname{argmin}_{1 \leq k \leq K} \|x_i - c_k\|$
3. For each of the clusters, the centroid is recalculated based on the data points contained in the cluster and minimizing the sum of squares within the cluster. That is,

$$WSS = \sum_{k=1}^K \sum_{C(i)=k} \|x_i - c_k\|^2$$

Go to step 2, until convergence is achieved.

K-Means



Alternatives for choosing the k initial centroids?

- Using the k first instances.
- Choosing randomly k instances.
- Taking randomly any partition into k clusters and calculating their centroids.
- Using the k-means++ method. D. Arthur and S. Vassilvitskii (2007). k-means++: the advantages of careful seeding. Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms. SODA's 07. pp. 1027-1035.
- Kmeans from scikit-learn uses this last method.

The k-means algorithm : Pro and cons

- **Pros:**
- The k-means algorithm is computationally fast.
- It works fine with data containing missing values.
- Complexity: $O(nkpi)$, n :number of instances, k =number of clusters, p :number of predictors, i :number of iterations. But for large datasets kmeans is basically $O(n)$.
- **Cons:**
- The optimization criterion it is not satisfied globally, only a local optimum is achieved.
- It is sensible to outliers.

kmeans in python

The class **sklearn.cluster** includes a function **KMeans** for performing the kmeans algorithm.

Example: K-means applied to diabetes(ignoring the class column)

```
kmeans = KMeans(n_clusters=2, random_state=0).fit(X)
clustlabels=kmeans.labels_
print clustlabels
#Choosing the two centroids at random
kmeans = KMeans(n_clusters=2, init='random', random_state=0).fit(X)
clustlabels1=kmeans.labels_
print clustlabels1
```

Partitioning around medoids (PAM)

Introduced by Kauffman and Rousseauw, 1987.

MEDOIDS, are data points representative of the clusters that are beeing formed

For a pre-specified number of clusters K, the PAM algorithm is based on the search **of the K MEDOIDS**,

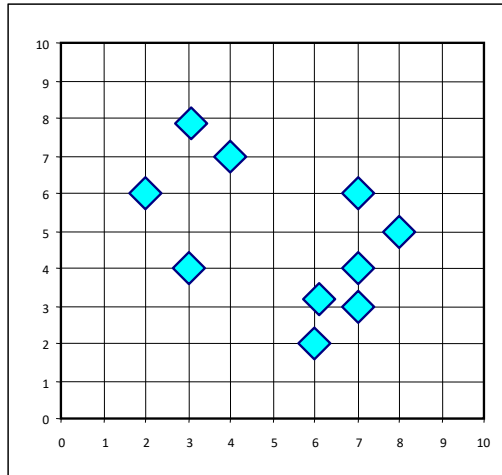
M = (**m**₁, . . . , **m**_K) of all the instances.

In order to find M, the sum of the distances of the data points to their closest medoid must be minimized. That is,

$$M^* = \arg \min_M \sum_i \min_k d(x_i, m_k)$$

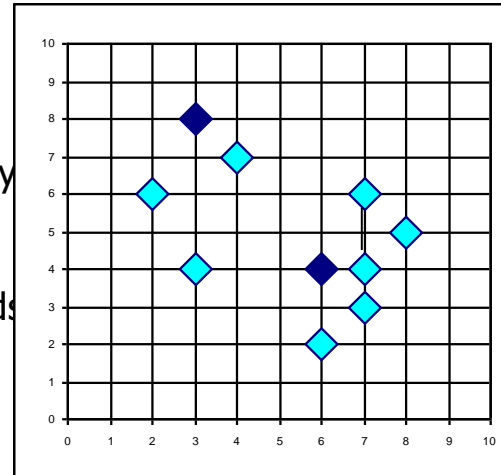
where d is a dissimilarity measure

k-Medoides



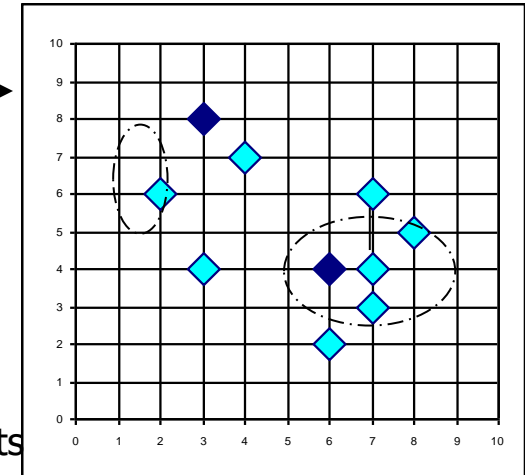
$K=2$

Arbitrarily
choose k
instances
as medoids



Total Cost = 26

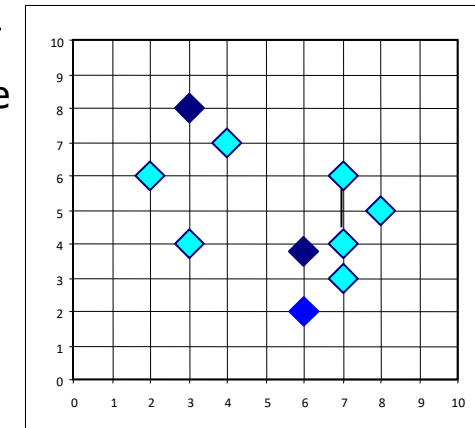
Assign
the data
points to
the cluster
where is
located its
closer
medoid.




Total Cost = 20

Choose randomly a
nonmedoid datapoint, O_{random}

Calculate the
cost total of
swapping



Swapping O
an O_{random}
If the Total
cost
decreases.



Continue the
loop until the
objective
function does
not change lo

PAM in Python

The scikit library does not include any function for performing PAM. There is a PyCluster library, where PAM is included

To apply **PAM** to the Diabetes dataset:

```
import Pycluster
```

```
Xd=np.matrix(Xd)
```

```
dist2=pairwise_distances(Xd)
```

```
clusterid,b,c=Pycluster.kmedoids(dist2)
```

PAM's complexity: $O(k(n-k)^2)$. It is too heavy for large datasets . There are two PAM versions CLARA and CLARANS for large datasets , but still the computation is not fast.

Self-organizing Maps, SOM (Kohonen, 1988)

SOM is a partitioning algorithm with the restriction that the clusters must be represented in a regular structure of low dimension such as a grid. SOM uses mainly one and two-dimensional grids.

The cluster that closer among them appear in adjacent cells of the grid. This means SOM maps the sample space into a space of lower dimension in which the similarity measure among the instances is measured by the close relationship of neighbors.

Each of the K clusters is represented by a prototype M_i , $i = 1, \dots, K$.

Self-organizing Maps, (SOM)

In this method the initial position of the data point in the sample space affects somehow its assignment to a cluster.

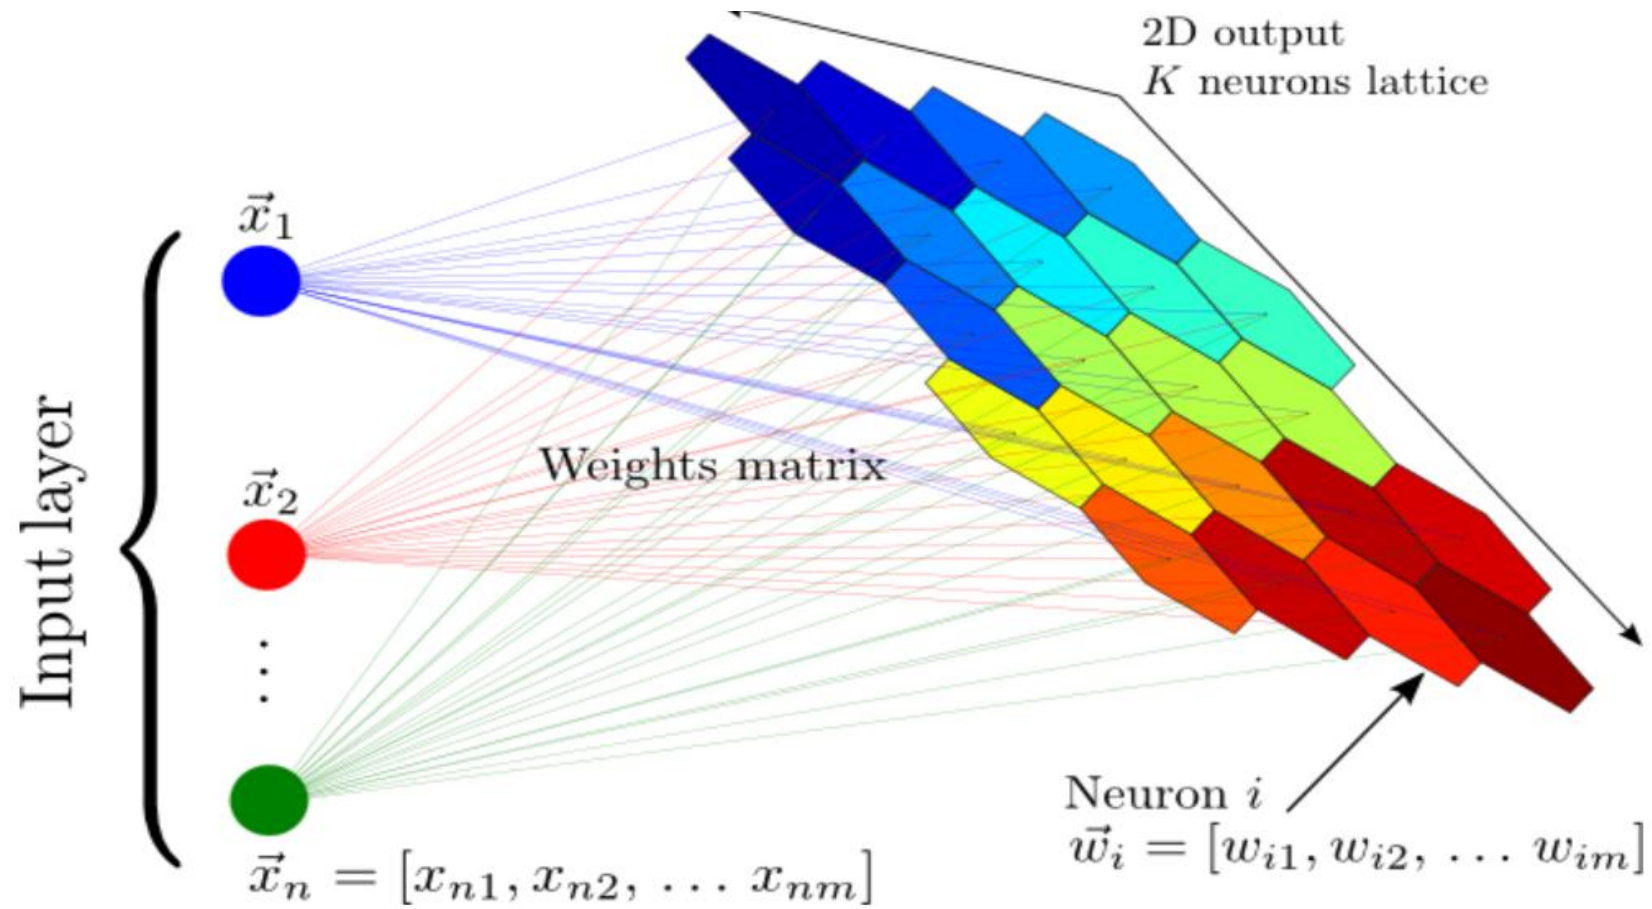
A SOM can be visualized and the datapoints that are similar among appear close in the plot.

The SOM algorithm is quite similar to k-means.

A SOM is trained as a neural network.

The SOM can also be considered as a dimensionality reduction technique. It is closed related to multidimensional scaling (process to find a set of data points in a multidimensional space that are similar among them).

Self-organizing Maps, (SOM)



SOM Algorithm considering a bidimensional rectangular grid

- Step 1**-Select the number of rows (q_1) and columns (q_2) to have in the grid. Therefore, the data will be clustered in $K=q_1q_2$ groups.
- Step 2**. Initialize the updating parameter (the learning rate in terms of neural nets) α ($\alpha=1$) and the radius of the grid ($r=2$)
- Step 3**. Initialize the prototype vectors $M_j, j \in (1, \dots, q_1) \times (1, \dots, q_2)$ by selecting randomly K instancias.
- Step 4**. for each data point x_i of the dataset do the following:
Identify the index j' of the prototype M_j that is closer to x_i .
Identify a set S of neighbors prototypes of $M_{j'}$. That is,
 $S = \{j: \text{distance}(j, j') < r\}$
- The distance can be Euclidean or any other one.
Update each element of S moving the corresponding prototype towards x_i :
 $M_j \leftarrow M_j + \alpha(x_i - M_j)$ for each $j \in S$
- Step 5**. Decrease the values of α and r in a predetermined amount and continue the iteration until reach convergence.

SOM in Python

The scikit learn library does not include the SOM algorithm. There is a function `somcluster` in the library `Pycluster`. Also there is a `SOMPY` library, but it runs only on Linux and MacOS

Other software: `Somtoolbox` in Matlab, `Cluster by Eissen` and `Genecluster` (MIT Center for Genome Research). R has at least 4 libraries for SOM. `Rapidminer` also performs som

SOM using PyCluster

```
import Pycluster
clusterid,celldata=Pycluster.somcluster(Xd,nxgrid=1,nygrid=2)
clus=pd.DataFrame(clusterid)
#comparing the clusters with the actual classes
pd.crosstab(clus.sum(axis=1),yd)
```

class	1	2
row_0		
0	186	141
1	314	127

II. Hierarchical algorithms for clustering

In this type of algorithms ordered sequences (Hierarchy) of clusters are generated. It can be merging smaller clusters or dividing larger clusters into smaller ones. The hierarchical structure is represented as a tree plot and it is called **Dendrogram**.

There are two types:

Agglomerative hierarchical algorithms (bottom-up, at the beginning each data point is a single).

Divisive hierarchical algorithms (top-down, at the beginning all the data points form one cluster).

The python libraries scikit-learn and scipy have functions to carry out only the Agglomerative hierarchical algorithm

Basic example

The Data

```
x=[3, 4, 6, 7, 8, 10, 15, 18]
```

```
y=[ 8, 9, 12, 17, 24, 20, 30, 28]
```

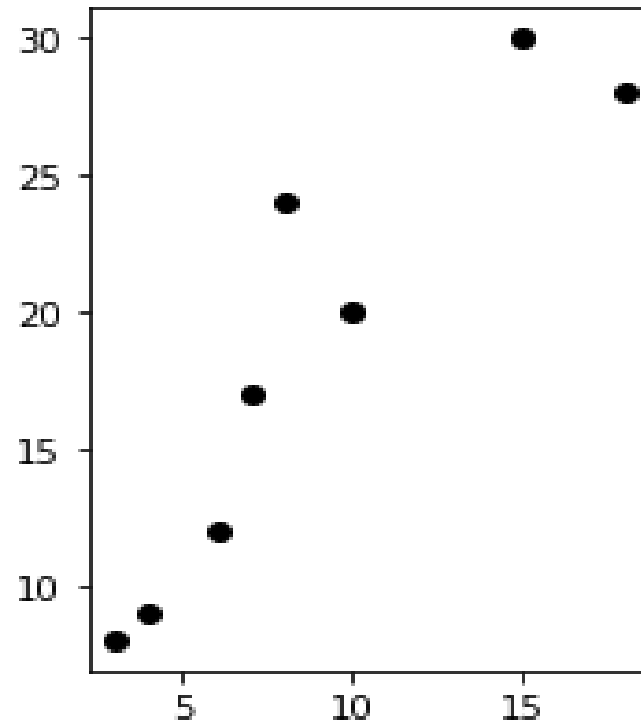
```
df=list(zip(x,y))
```

```
df=pd.DataFrame(df,columns=['x','y'])
```

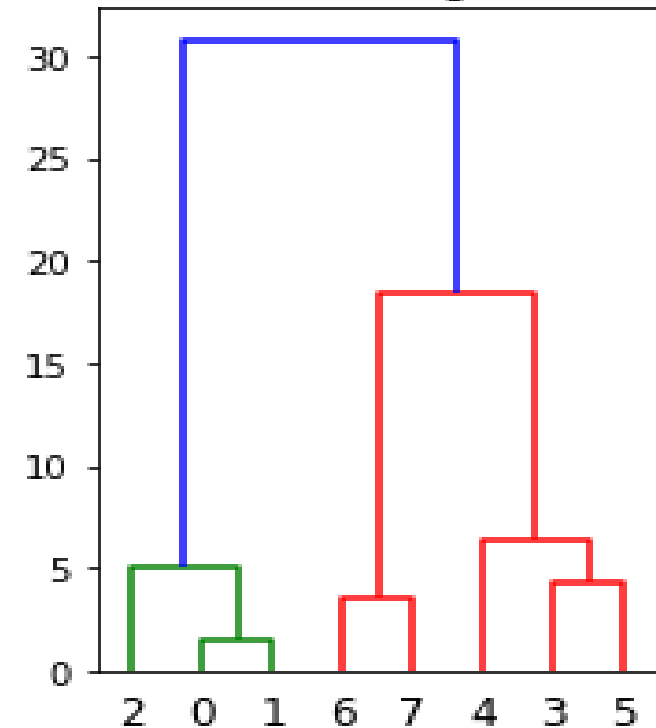
```
dist2=pairwise_distances(df)
```

```
Z = hierarchy.linkage(df,'ward')
```

The data



the dendrogram



Distance between clusters=Linkage functions

After choosing the dissimilarity measure, the data matrix $n \times p$ (n data points, p features) is transformed in distance matrix or dissimilarity matrix $D = (d_{ij})$ of order $n \times n$ for the data points to be grouped.

Based on this matrix, a distance measure between two clusters S and T must be computed. The following are some of this distance measures:

Single Linkage: $\delta(S,T) = \min_{\{x \in S, y \in T\}} d(x,y)$

Complete Linkage: $\delta(S,T) = \max_{\{x \in S, y \in T\}} d(x,y)$

Average Linkage:

$$\delta(S,T) = \frac{1}{|S| + |T|} \sum_{x \in S, y \in T} d(x,y)$$

where $|S|$, and $|T|$ represent the cardinality of S and T , respectively.

Centroid Linkage : $\delta(S,T) = d(\bar{x}, \bar{y})$

where \bar{x} and \bar{y} represent the centroids of S and T respectively.

Distance between clusters=Linkage functions

Median Linkage: $\delta(S,T)=\text{median}_{\{x \in S, y \in T\}} d(x,y)$

Mc Quitty Linkage:

$$\partial(S,T) = \frac{1}{|S|+|T|} \left[\sum_{x \in S} d(x, \bar{x}) + \sum_{y \in T} d(y, \bar{y}) \right]$$


Ward Linkage:

$$\partial(S,T) = \sum_{x \in S} d^2(x, \bar{x}) + \sum_{y \in T} d^2(y, \bar{y})$$

Ward linkage merges the pair of clusters that gives the smallest variance in the merged group.

Scikit-learn considers all of the above linkage functions.

Distance between clusters=Linkage functions



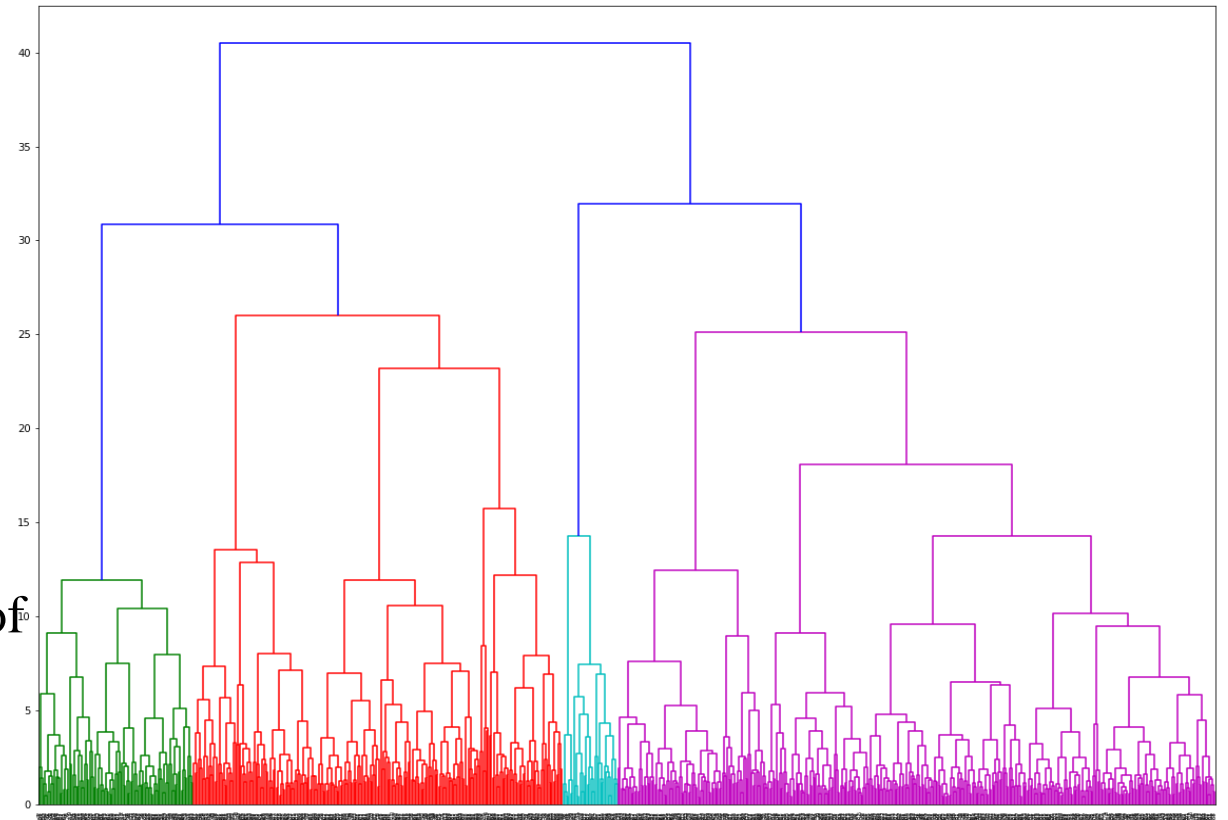
Hierarchical Algorithm applied to Diabetes

Since Diabetes is supervised, we will apply clustering only to the predictor matrix, X . First, we must apply a distance measure X (frequently the data is standardized) usually the euclidean distance.

Second, we apply the hierarchical algorithm to this distance matrix using a linkage function

```
Z = hierarchy.linkage(X,'ward')
```

The algorithm does not require the number of clusters as a parameter. The clustering assignment is done later on.

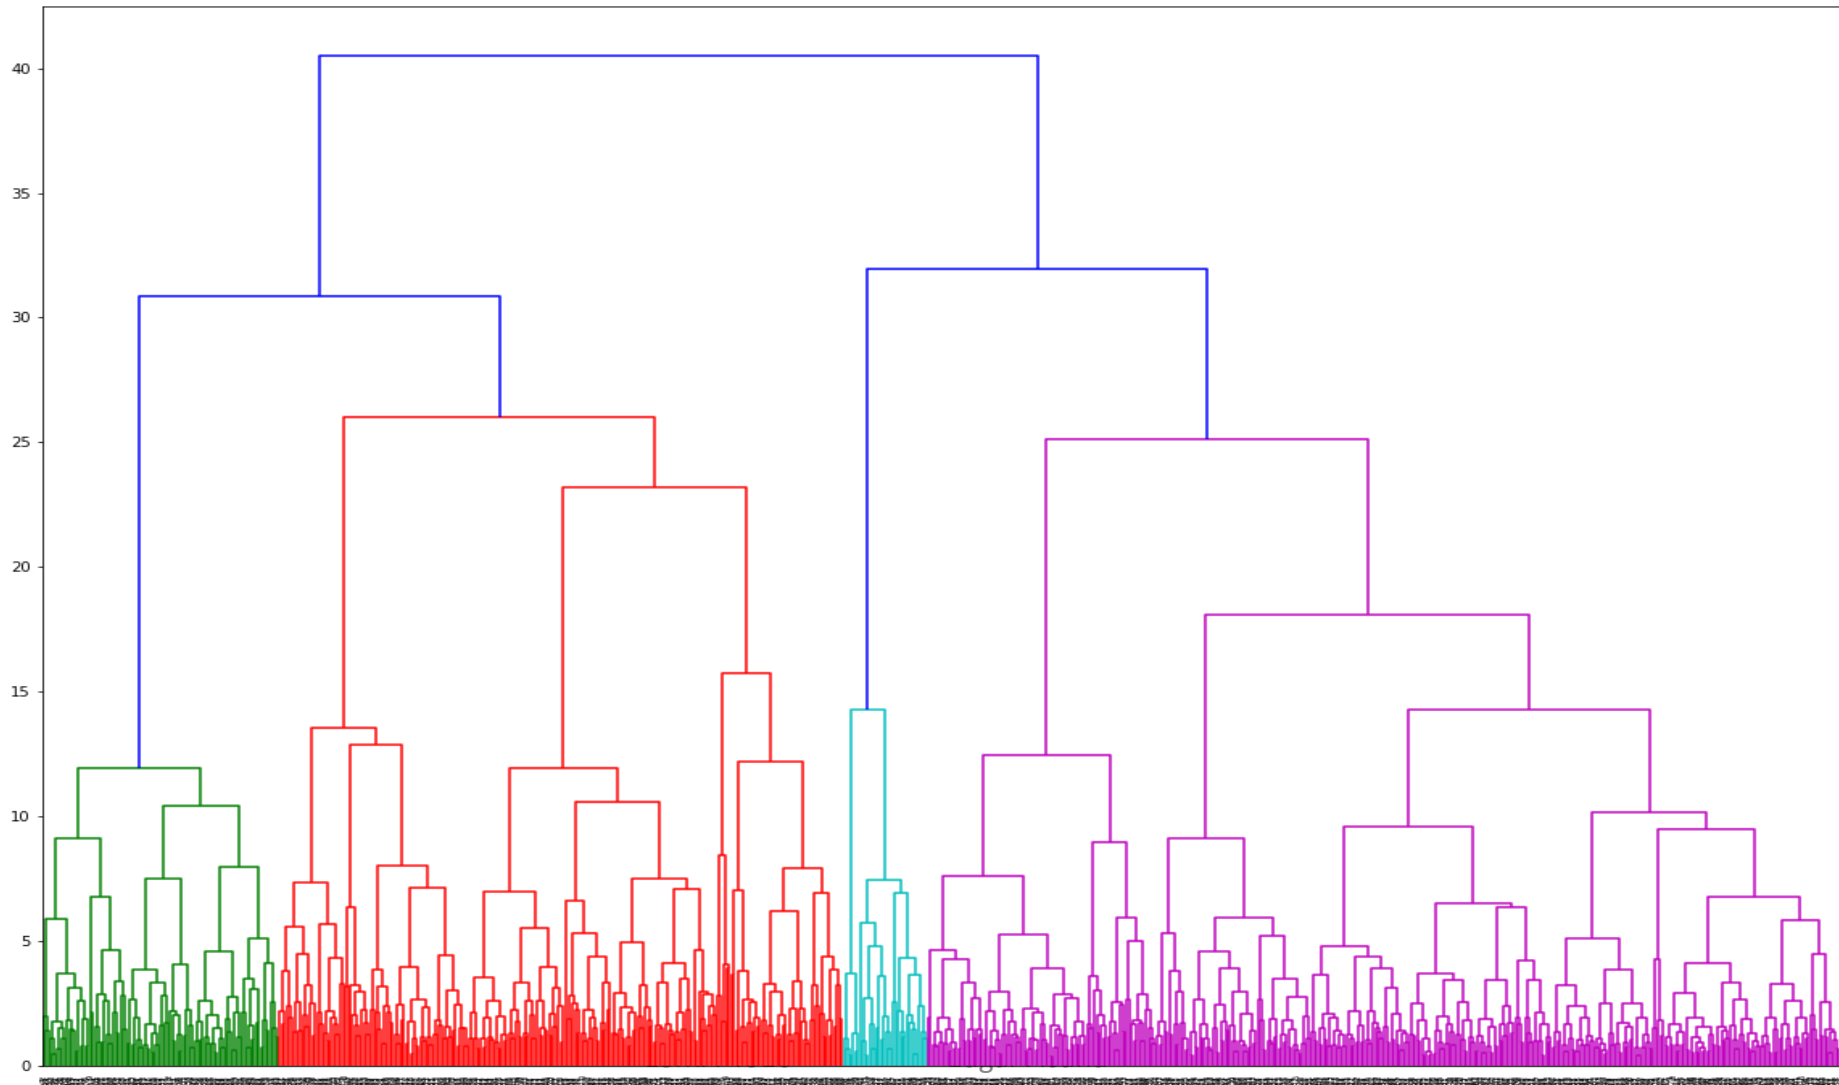


Dendrograms

The dendrograms are easy to understand but they can lead to wrong conclusions for the following reasons:

- 1) The corresponding dendrogram to a hierarchical clustering is not unique, since for each merge of clusters one needs to specify which subtree goes to right and which to the left. Most of the programs order the trees in such way that clusters including more data points go to the left.
- 2) The hierarchical structure of the dendrogram does not represent with certainty the true distances between the distinct data points of the dataset.

Dendrogram for Diabetes (Ward Linkage)



Dendrogram

The cophenetic correlation coefficient can be used as a measure of how well the hierarchical structure of the dendrogram represents the actual distances. It is defined as the correlation between the $n(n - 1)/2$ pairs of dissimilarities and their cophenetic distances in the dendrogram.

The library `scipy` allow us to compute the cophenetic correlation coefficient.

For the example,

```
from scipy.cluster.hierarchy import dendrogram, linkage
from scipy.cluster.hierarchy import cophenet
from scipy.spatial.distance import pdist
```

```
Z = linkage(df, 'ward')
c, coph_dists = cophenet(Z, pdist(df))
print c
```

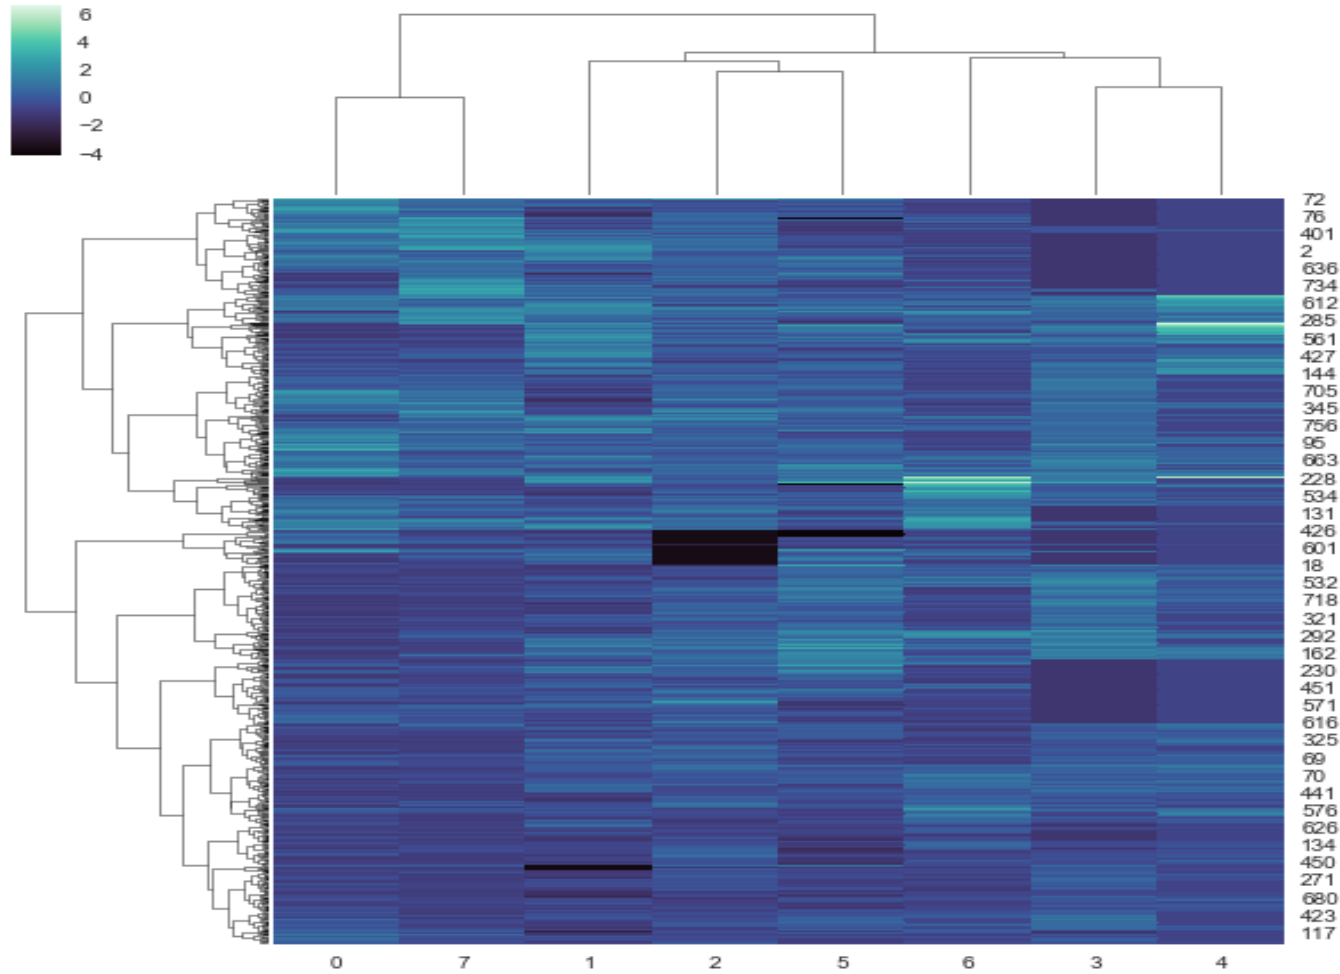
```
0.708252468327
```


Heatmaps.

Heatmaps are graphs that show simultaneously the cluster of rows and columns.

The visualization library seaborn has nice heatmaps

Heatmap for Diabetes



Hierarchical Agglomerative Clustering for Diabetes

```
from sklearn.cluster import AgglomerativeClustering
from sklearn.metrics import pairwise_distances
n_clusters = 2
model = AgglomerativeClustering(n_clusters=n_clusters,linkage="average")
model.fit(X)
clustlabels=model.fit_predict(X)
unique_elements, counts_elements = np.unique(clustlabels, return_counts=True)
print("Size of the two clusters")
print(np.asarray((unique_elements, counts_elements)))
Size of the two clusters
[[ 0 1]
 [761 7]]
The average linkage is very affected by outliers
```

Hierarchical Agglomerative Clustering for Diabetes

```
model = AgglomerativeClustering(n_clusters=n_clusters,linkage="ward")
model.fit(X)
clustlabels=model.fit_predict(X)
unique_elements, counts_elements = np.unique(clustlabels, return_counts=True)
print("Size of the two clusters")
print(np.asarray((unique_elements, counts_elements)))
Size of gthe two clusters
[[ 0 1]
 [426 342]]
```

We get much better results using the Ward Linkage

Partitioning Algorithms versus Hierarchical Algorithms

The partitioning methods have the advantage that a optimality criterion is satisfied at least approximately. They are fast and can be parallelized. The main disadvantage is that they need a number of clusters to start working.

On other hand, the hierarchical algorithms have a time complexity of at least $O(n^2)$.

The agglomerative are much faster to compute than the divisive.

Other disadvantage of these methods is that they suffer of the nesting problema. The tree structure is rigid, and it is not easy make corrections to the clustering process.